

Transformers & Large Language Models (LLMs)

From Concepts to Practice - ML Course Day 4

Presented by

Mahamad Salah

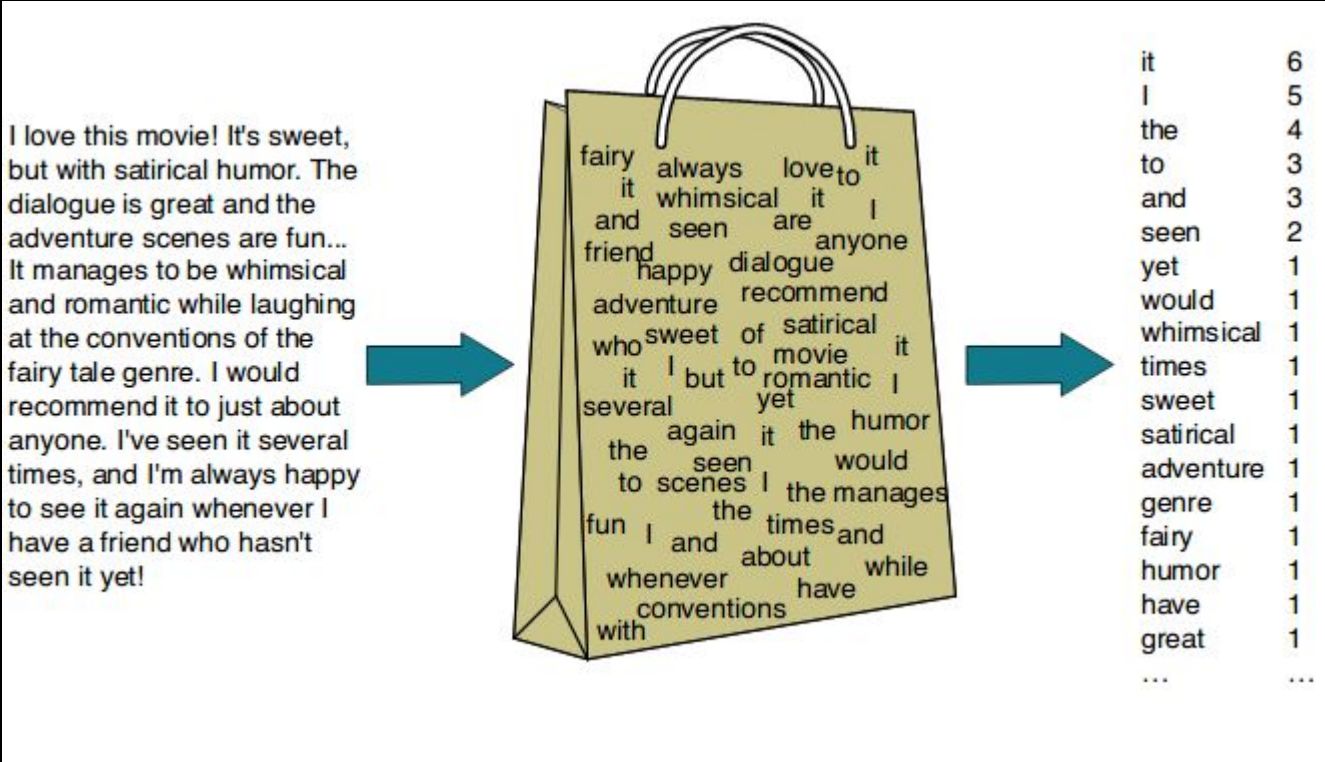
Anton Popov

Content



- 1 Attention is All You Need! >
- 2 The Anatomy of a Transformer >
- 3 An Intro to LLMs >
- 4 HuggingFace & Open Source Models >
- 5 NLP Problems: A Demo >
- 6 Transformers Beyond Text >

Attention is All You Need



TF*IDF

**TF*IDF = TERM FREQUENCY * INVERSE
DOCUMENT FREQUENCY**

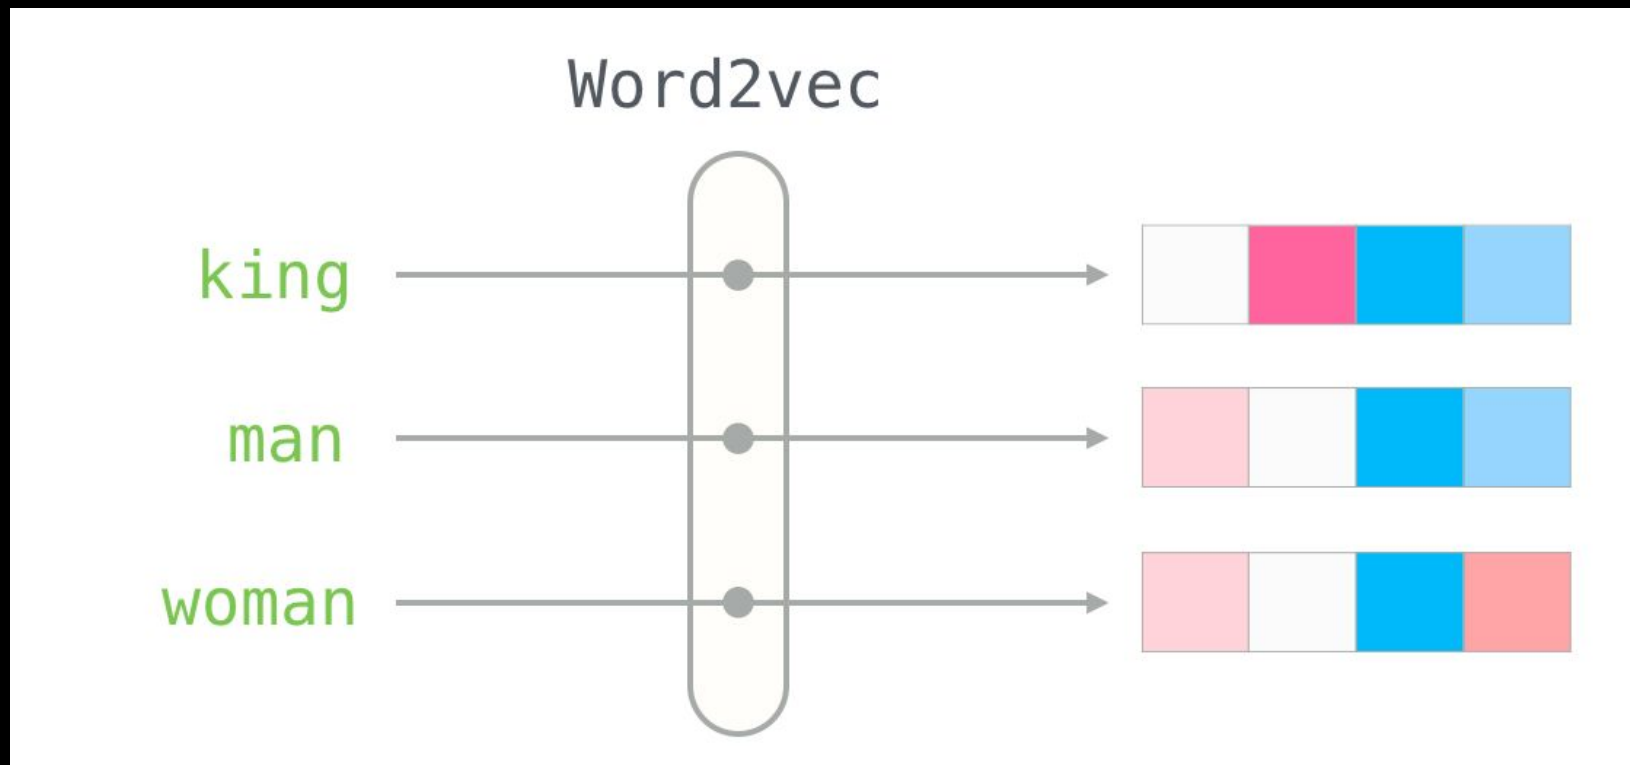
TERM FREQUENCY=

THE AMOUNT OF TIMES A
TERM APPEARS IN A
DOCUMENT



X

INVERSE DOCUMENT
FREQUENCY=

A MEASURE OF WHETHER A
TERM IS RARE OR COMMON
IN A COLLECTION OF
DOCUMENTS.



Attention is All You Need



Cosine similarity:

$$\cos \theta = \frac{\vec{a} \cdot \vec{b}}{\|\vec{a}\| \cdot \|\vec{b}\|}$$

EN2	EN1
Entrez:2020	Entrez:2020


$$\cos \theta = 0.99$$

Nausea	Headache
MESH:D009325	MESH:D006261


$$\cos \theta = 0.976$$

Cyclophosphamide	Neutropenia
MESH:D003520	MESH:D009503


$$\cos \theta = 0.819$$




Attention is All You Need



RNN



LSTM



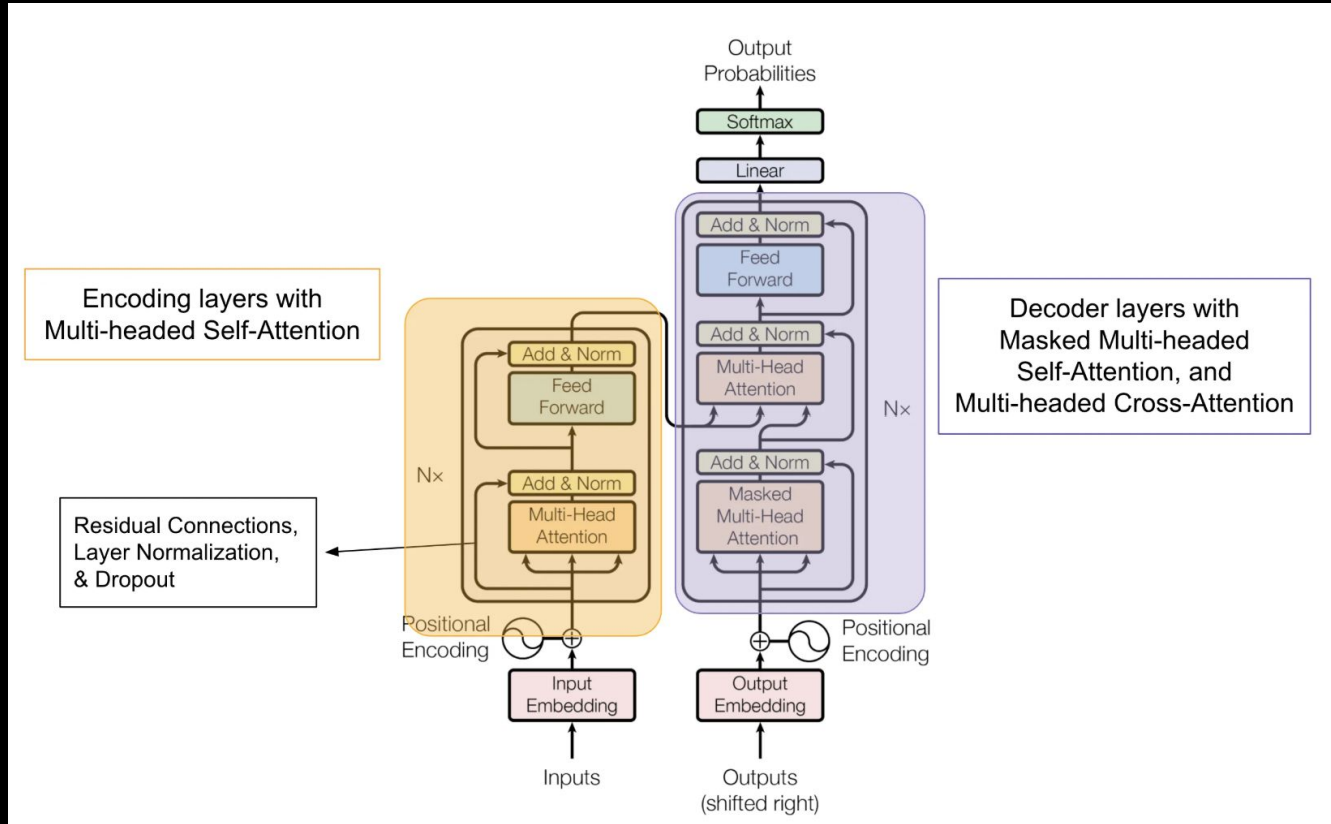
GRU

Attention is All You Need

Major Challenges:

- Context Dependence
- Order Awareness
- Long-range dependencies
- Vanishing gradients

The Anatomy of a Transformer



The Anatomy of a Transformer

1) This is our input sentence*

Thinking
Machines



2) We embed each word*




* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one




3) Split into 8 heads. We multiply X or R with weight matrices



...



4) Calculate attention using the resulting $Q/K/V$ matrices



...



5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer




...

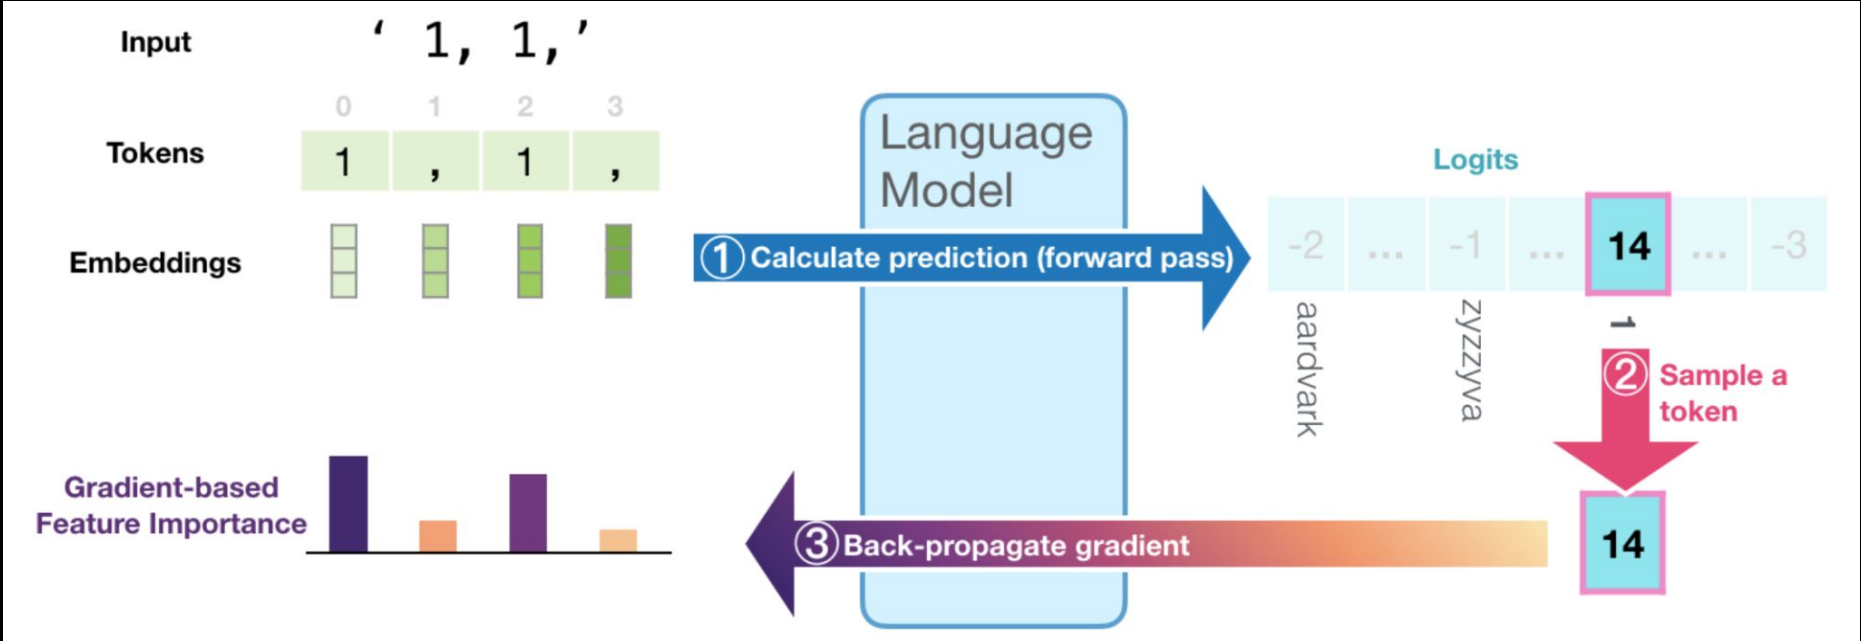


W^O



Z







Hugging Face

- Sentence Classification
- Named Entity Recognition
- Text Generation
- Questions & Answers

Beyond Text: Vision Transformers

